

rtsm-code-1

April 15, 2024

1 RTSM Project

```
[1]: import numpy as np
from scipy.stats import t, chi2,f
import matplotlib.pyplot as plt
import seaborn as sns
from numpy import genfromtxt
from sklearn.preprocessing import PolynomialFeatures
import cvxpy as cp
import random
import math
import pandas as pd
import itertools
from sklearn.model_selection import train_test_split
```

2 Dataset Loading and EDA

```
[2]: df = pd.read_csv('Concrete_Data.csv')
```

```
[3]: df.head()
```

```
[3]:
```

	Cement	Blast Furnace Slag	Fly Ash	Water	Superplasticizer \
0	540.0	0.0	0.0	162.0	2.5
1	540.0	0.0	0.0	162.0	2.5
2	332.5	142.5	0.0	228.0	0.0
3	332.5	142.5	0.0	228.0	0.0
4	198.6	132.4	0.0	192.0	0.0

	Coarse Aggregate	Fine Aggregate	Age	Concrete compressive strength
0	1040.0	676.0	28	79.99
1	1055.0	676.0	28	61.89
2	932.0	594.0	270	40.27
3	932.0	594.0	365	41.05
4	978.4	825.5	360	44.30

```
[4]: df.describe()
```

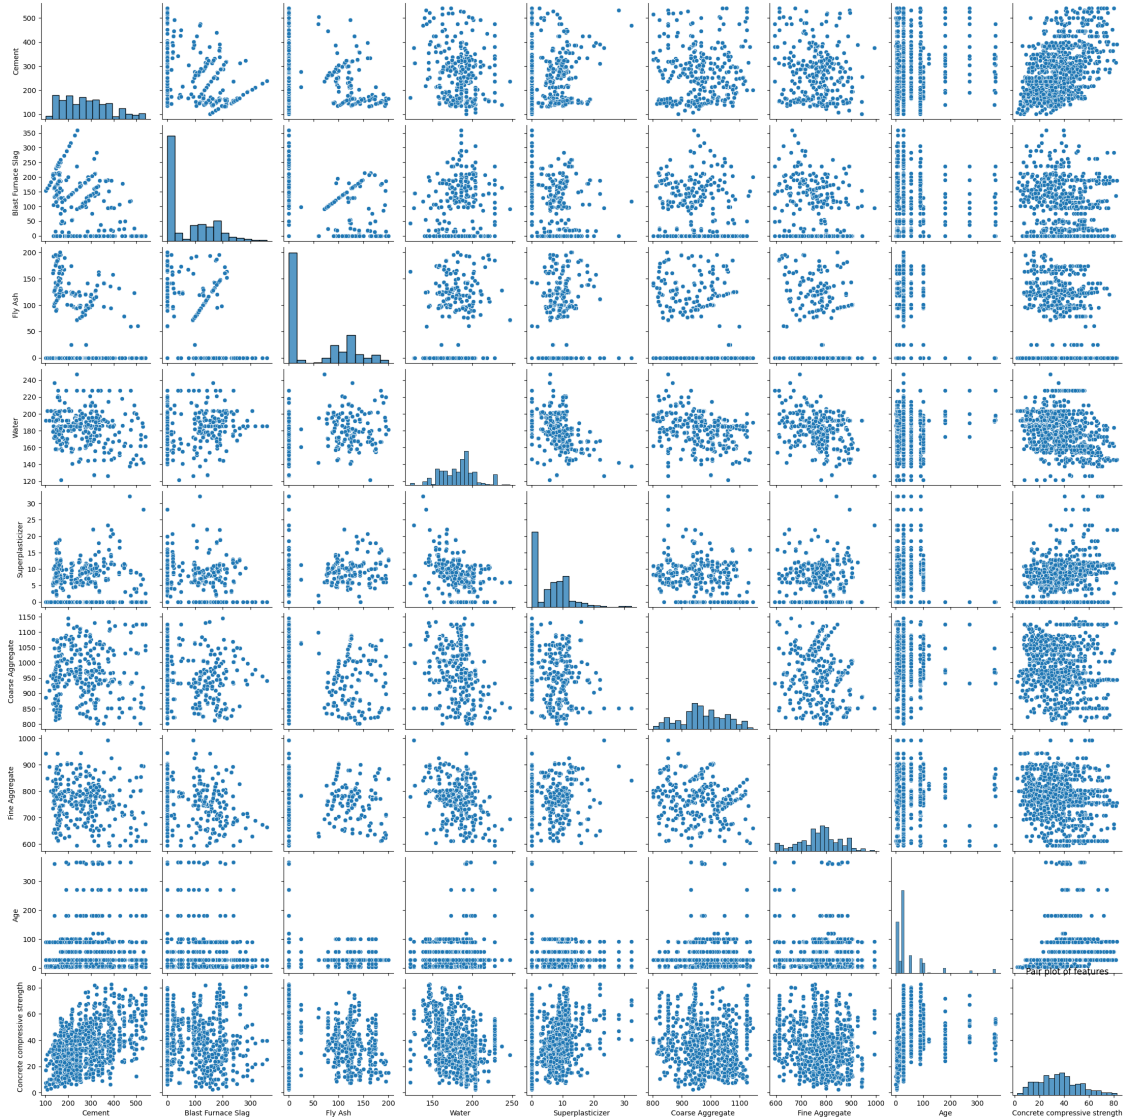
```
[4]:
```

	Cement	Blast Furnace Slag	Fly Ash	Water	\
count	1030.000000	1030.000000	1030.000000	1030.000000	
mean	281.166408	73.894854	54.187379	181.564854	
std	104.507710	86.279340	63.995962	21.355663	
min	102.000000	0.000000	0.000000	121.800000	
25%	192.375000	0.000000	0.000000	164.900000	
50%	272.900000	22.000000	0.000000	185.000000	
75%	350.000000	142.950000	118.300000	192.000000	
max	540.000000	359.400000	200.100000	247.000000	

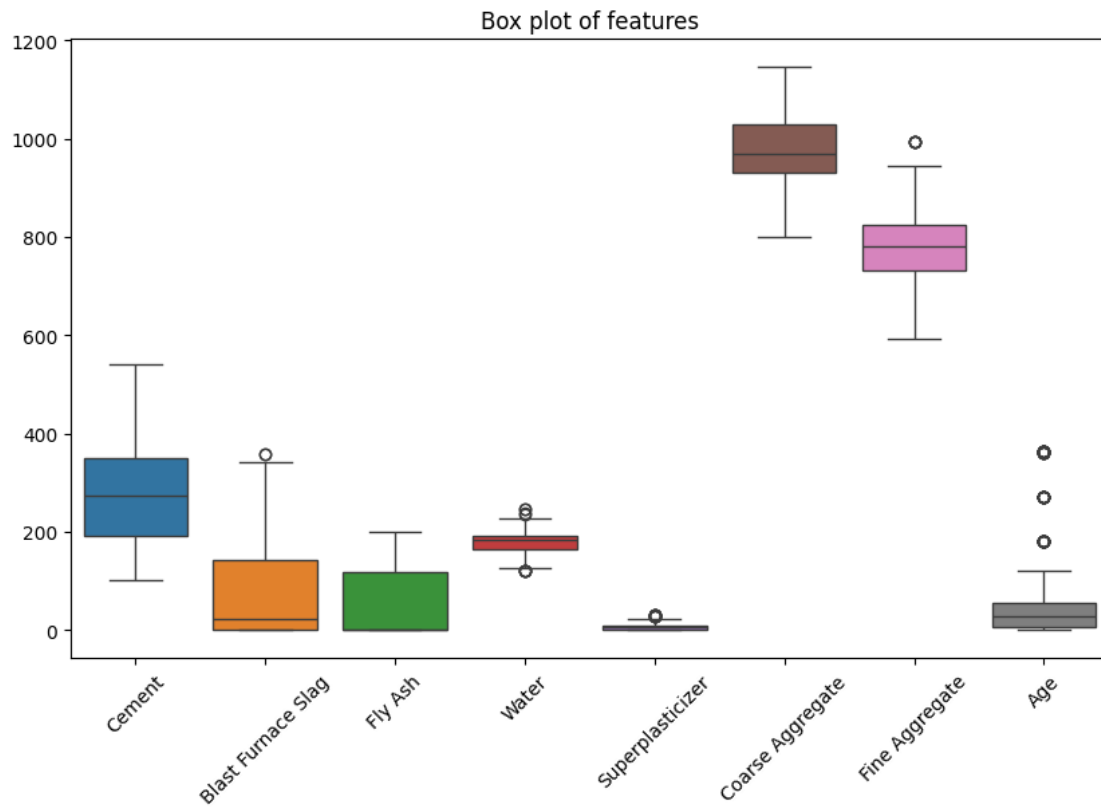
	Superplasticizer	Coarse Aggregate	Fine Aggregate	Age	\
count	1030.000000	1030.000000	1030.000000	1030.000000	
mean	6.203204	972.918932	773.579515	45.662136	
std	5.973035	77.753954	80.175801	63.169912	
min	0.000000	801.000000	594.000000	1.000000	
25%	0.000000	932.000000	730.950000	7.000000	
50%	6.300000	968.000000	779.500000	28.000000	
75%	10.200000	1029.400000	824.000000	56.000000	
max	32.200000	1145.000000	992.600000	365.000000	

	Concrete compressive strength
count	1030.000000
mean	35.817961
std	16.705742
min	2.330000
25%	23.710000
50%	34.445000
75%	46.135000
max	82.600000

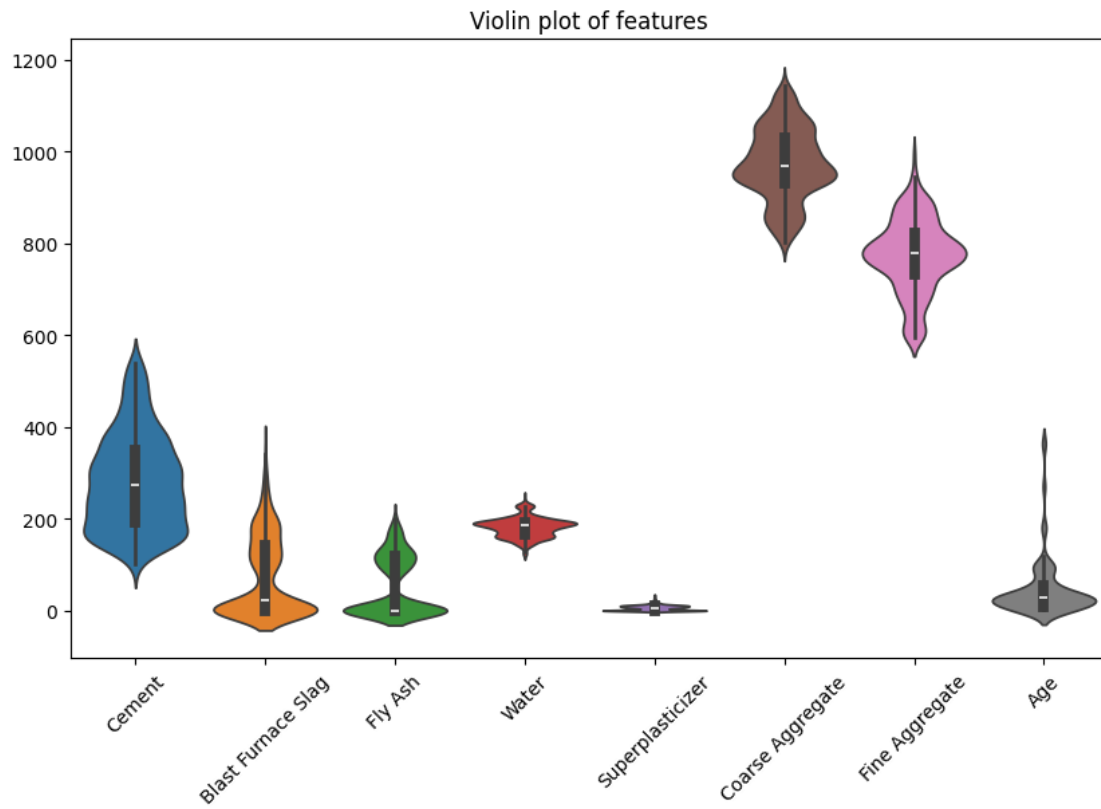
```
[5]: sns.pairplot(df)
plt.title('Pair plot of features')
plt.show()
```




```
[6]: plt.figure(figsize=(10, 6))
sns.boxplot(data=df.drop('Concrete compressive strength', axis=1))
plt.title('Box plot of features')
plt.xticks(rotation=45)
plt.show()
```



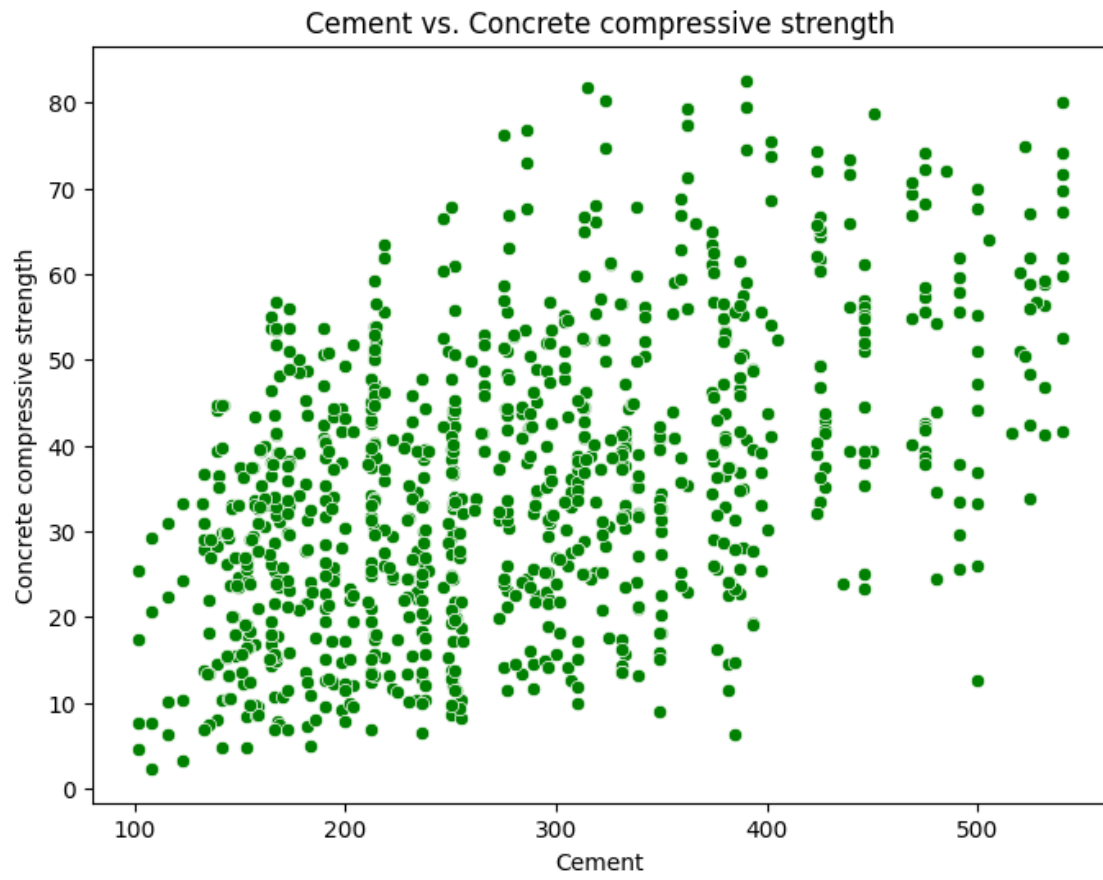
```
[7]: plt.figure(figsize=(10, 6))
sns.violinplot(data=df.drop('Concrete compressive strength', axis=1))
plt.title('Violin plot of features')
plt.xticks(rotation=45)
plt.show()
```

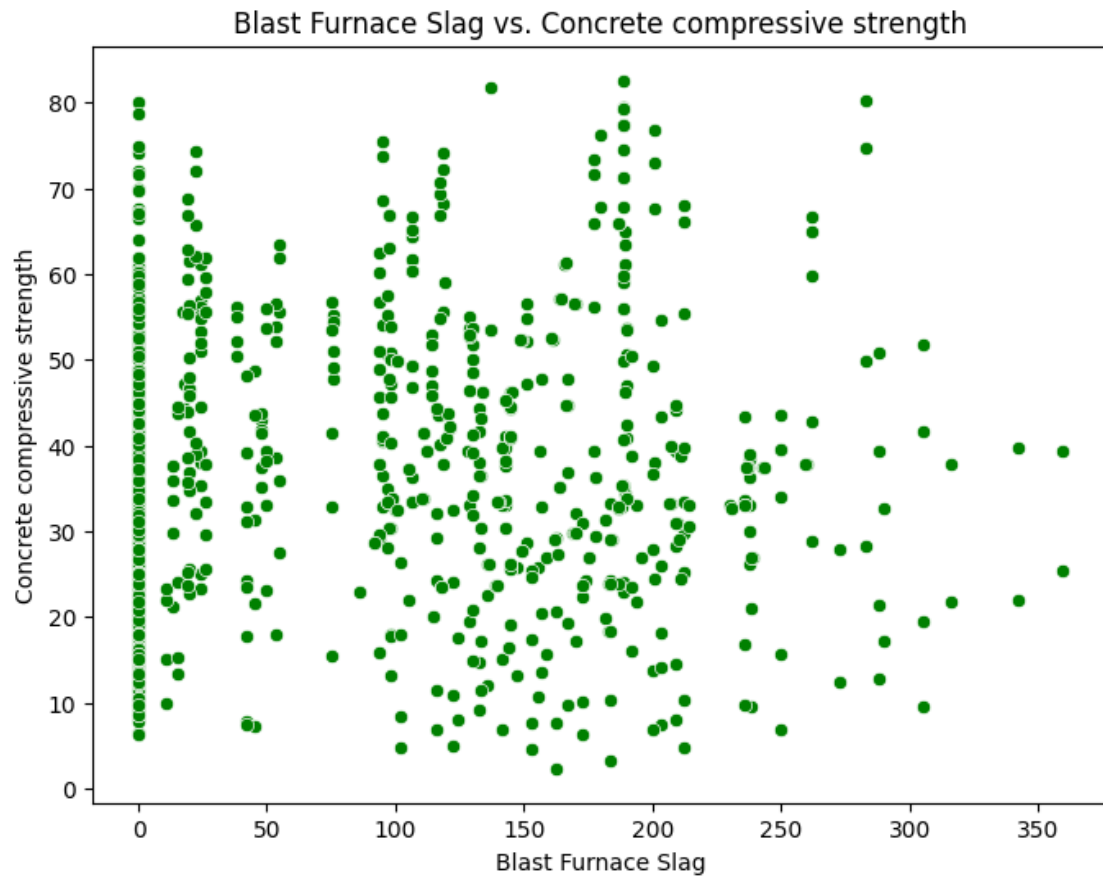


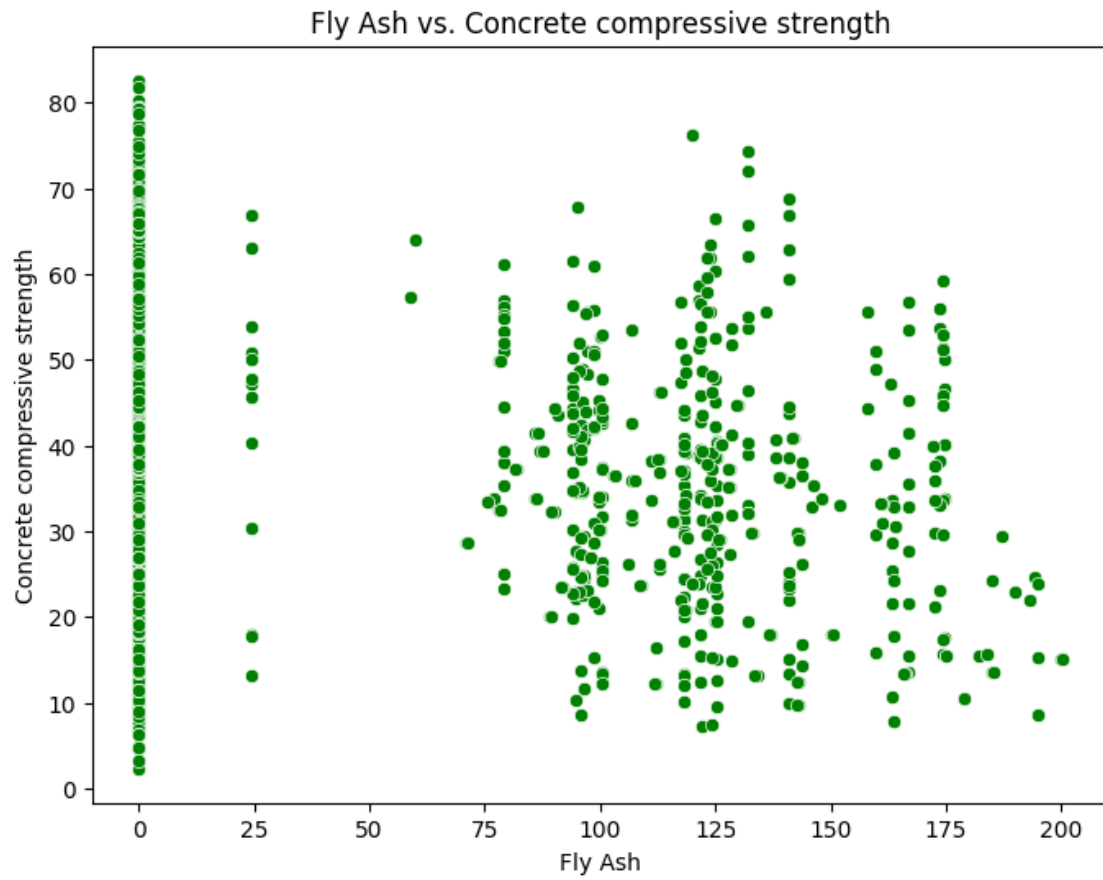
```
[8]: plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation heatmap')
plt.show()
```




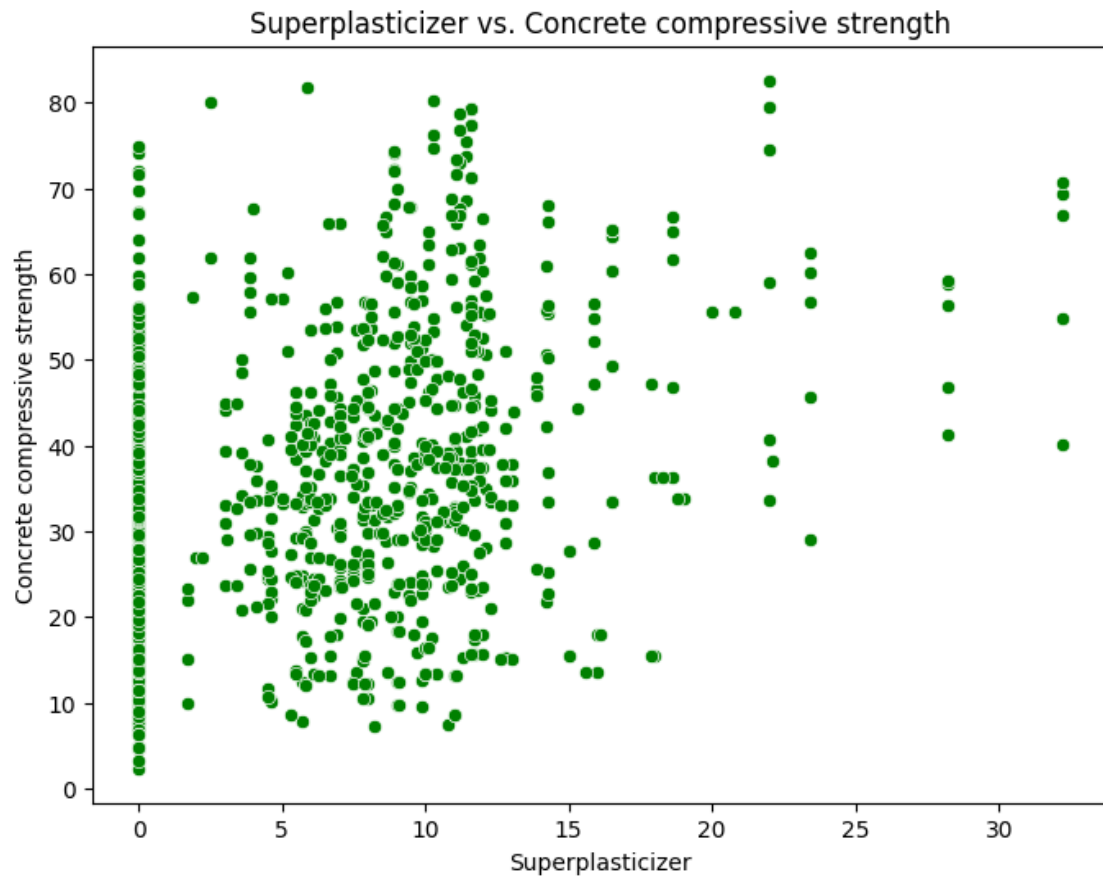
```
[9]: for column in df.columns:
    if column != 'Concrete compressive strength':
        plt.figure(figsize=(8, 6))
        sns.scatterplot(data=df, x=column, y='Concrete compressive strength',
            color='green')
        plt.title(f'{column} vs. Concrete compressive strength')
        plt.xlabel(column)
        plt.ylabel('Concrete compressive strength')
        plt.show()
```




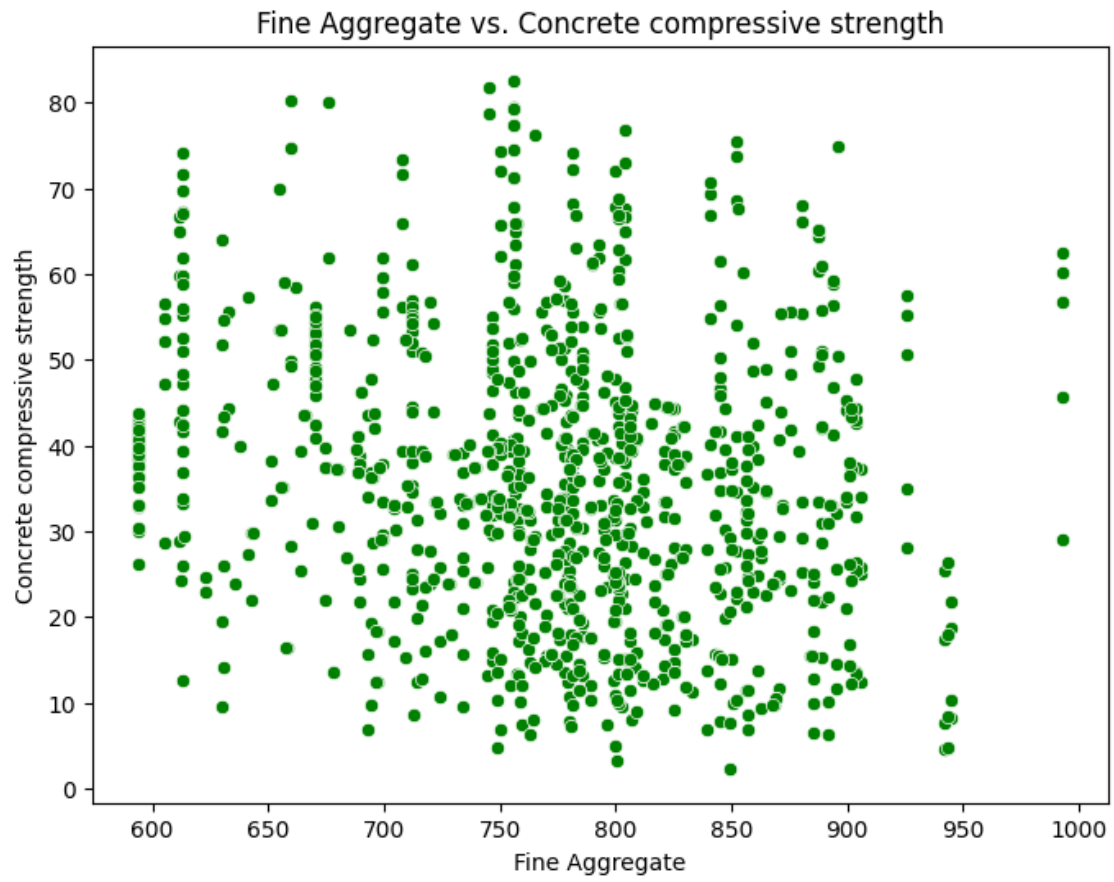


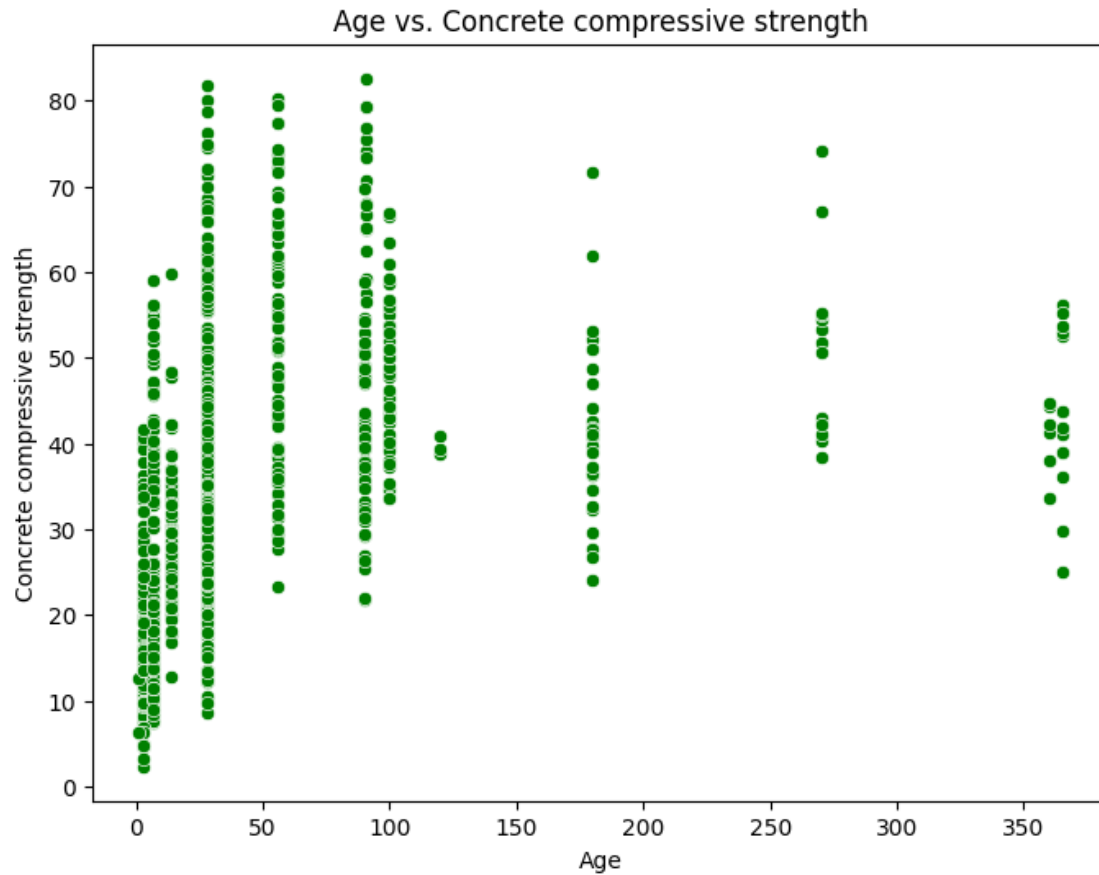












3 Multiple Regression

```
[10]: import numpy as np
from scipy.stats import t, chi2, f

def fit(X, Y):
    beta = np.linalg.inv(X.T @ X) @ X.T @ Y
    return beta

def analyse_model(X, Y, beta, X_test = None, Y_test = None,
    print_terms=False, test_check = True):
    k = X.shape[1] - 1
    n = Y.shape[0]

    Y_est = X @ beta
    e = Y - Y_est
    sigma_square_est = (e.T @ e) / (n - k - 1)
```

```

significance = []

# Testing significance of coefficients
C = np.linalg.inv(X.T @ X)
alpha = 0.05
test_stat_beta = np.zeros(k + 1)

for i in range(k + 1):
    test_stat_beta[i] = beta[i] / (np.sqrt(sigma_square_est) * np.sqrt(C[i,
↪i]))

significant_count = 0
non_significant_count = 0
for i in range(k + 1):
    if np.abs(test_stat_beta[i]) > t.ppf(1 - alpha / 2, n - k - 1, loc=0,
↪scale=1):
        significance.append(1)
        significant_count += 1
        if print_terms:
            print("beta", i, " is significant")
    else:
        significance.append(0)
        non_significant_count += 1
        if print_terms:
            print("beta", i, " is not significant")

# Confidence intervals
l_interval_beta = np.zeros(k + 1)
h_interval_beta = np.zeros(k + 1)

for i in range(k + 1):
    l_interval_beta[i] = beta[i] - np.sqrt(sigma_square_est) * t.ppf(1 -
↪alpha / 2, n - k - 1, loc=0, scale=1) * np.sqrt(C[i, i])
    h_interval_beta[i] = beta[i] + np.sqrt(sigma_square_est) * t.ppf(1 -
↪alpha / 2, n - k - 1, loc=0, scale=1) * np.sqrt(C[i, i])
    if print_terms:
        print("95% confidence interval for beta", i, ":", l_interval_beta[i],
↪",", h_interval_beta[i])

    l_interval_sigma_square = (n - k - 1) * sigma_square_est / chi2.ppf(1 -
↪alpha / 2, df=n - k - 1)
    h_interval_sigma_square = (n - k - 1) * sigma_square_est / chi2.ppf(alpha /
↪2, df=n - k - 1)
    if print_terms:
        print("95% confidence interval for sigma square :",
↪l_interval_sigma_square, ",", h_interval_sigma_square)

```

```

# ANOVA + coefficient of determination
SSError = e.T @ e
SSTotal = (Y - np.mean(Y)).T @ (Y - np.mean(Y))
SSReg = SSTotal - SSError
test_stat_anova = (SSReg * (n - k - 1)) / (SSError * k)

if ((test_stat_anova < f.ppf(1 - alpha / 2, dfn=k, dfd=n - k - 1, loc=0,
↪scale=1)) and (
    test_stat_anova > f.ppf(alpha / 2, dfn=k, dfd=n - k - 1, loc=0,
↪scale=1))) or non_significant_count==n:
    model_significance = "insignificant"
else:
    model_significance = "significant"

print("The Sum of Squared Errors for the model ( SSE ) :",SSError)
print("The Mean Squared error for the model ( MSE ) :",SSError/(n-1))
print("The Sum of Squares Regression for the model ( SSR ) :",SSReg)
print("The Total error for the model ( SST ) : ",SSTotal)
print("The Model is ",model_significance)

R_2 = 1 - (SSError / SSTotal)
R_2_adjusted = 1 - (SSError / SSTotal) * ((n - 1) / (n - k - 1))

# Printing results
print("Multiple Regression Model Summary:")
print("-----")
print("R-squared           : {:.4f}".format(R_2))
print("Adjusted R-squared   : {:.4f}".format(R_2_adjusted))
print("Number of significant terms : {}".format(significant_count))
print("Number of non-significant terms: {}".format(non_significant_count))
print("Model significance    : {}".format(model_significance))

if test_check:
    # Test set
    Y_test_pred= X_test @ beta
    SSE = np.sum((Y_test_pred - Y_test)**2)
    print("SSE on test set:",SSE)
    print("MSE on test set:",SSE/X_test.shape[1])

    print("-----")

return R_2, R_2_adjusted,significance

```

```

[11]: import numpy as np
import matplotlib.pyplot as plt

```



```

from scipy.stats import chi2

def plot_residual(X, Y, beta, k):
    error = Y - X @ beta
    H = X @ np.linalg.inv(X.T @ X) @ X.T
    MSEError = error.T @ error / (X.shape[0] - k - 1)
    d = error / np.sqrt(MSEError)
    r = np.array([error[i] / np.sqrt(MSEError * (1 - H[i, i])) for i in range(Y.
↪shape[0])])

    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

    # Scatter plot of Y vs. r
    ax1.scatter(X@beta, r, color='blue', alpha=0.5)
    ax1.set_xlabel('Y (predicted)')
    ax1.set_ylabel('Studentized Residuals (r)')
    ax1.set_title('Residual Plot')
    ax1.grid(True)

    # Histogram of studentized residuals
    ax2.hist(r, bins=20, density=True, alpha=0.6, color='green')
    ax2.set_xlabel('Studentized Residuals (r)')
    ax2.set_ylabel('Frequency')
    ax2.set_title('Histogram of Studentized Residuals')
    ax2.grid(True)

    # Calculate chi-square statistic and p-value
    observed_counts, _ = np.histogram(r, bins=20, density=True)
    expected_counts = np.mean(observed_counts) * np.ones_like(observed_counts)
    chi2_statistic = np.sum((observed_counts - expected_counts) ** 2 /
↪expected_counts)

    # Calculate degrees of freedom
    df = len(observed_counts) - 1

    # Calculate p-value
    p_value = chi2.sf(chi2_statistic, df)

    # Print chi-square test result
    print("Chi-square statistic:", chi2_statistic)
    print("P-value:", p_value)

    plt.tight_layout()
    plt.show()

```

4 Fitting data

```
[12]: X_data = np.array(df.drop(['Concrete compressive strength'],axis = 1))

Y_data = np.array(df['Concrete compressive strength'])

from sklearn.preprocessing import StandardScaler, MinMaxScaler
scaler_X = StandardScaler()
X_scaled = scaler_X.fit_transform(X_data)

# Scale labels (Y_data)
scaler_Y = StandardScaler()
Y_scaled = scaler_Y.fit_transform(Y_data.reshape(-1, 1)).reshape(-1)

X = X_scaled.copy()
Y = Y_scaled.copy()
n = np.size(Y)
X = np.concatenate((np.ones((n,1)),X),axis=1)

[13]: from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
↳random_state=42)
beta = fit(X_train,Y_train)
for i in range(X.shape[1]):
    print(f"The value of beta[i] is {beta[i]}")

analyse_model(X_train,Y_train,beta,X_test, Y_test,True)
plot_residual(X_train, Y_train, beta, X.shape[1]-1)
```

```
The value of beta0 is -0.00022021634322658828
The value of beta1 is 0.7443133159110672
The value of beta2 is 0.5604716950282813
The value of beta3 is 0.31472507060666927
The value of beta4 is -0.17335909258652094
The value of beta5 is 0.11063442201274962
The value of beta6 is 0.0829256551997025
The value of beta7 is 0.11548187137899125
The value of beta8 is 0.4403094364220793
beta 0  is not significant
beta 1  is significant
beta 2  is significant
beta 3  is significant
beta 4  is significant
beta 5  is significant
beta 6  is not significant
beta 7  is significant
```

beta 8 is significant

95% confidence interval for beta 0 : -0.043574477891147524 , 0.04313404520469435

95% confidence interval for beta 1 : 0.6290990502749305 , 0.8595275815472039

95% confidence interval for beta 2 : 0.44572749558127067 , 0.675215894475292

95% confidence interval for beta 3 : 0.20836671556076408 , 0.42108342565257445

95% confidence interval for beta 4 : -0.28409053769322345 , -0.06262764747981844

95% confidence interval for beta 5 : 0.03743013959424246 , 0.18383870443125677

95% confidence interval for beta 6 : -0.012066017480366342 , 0.17791732787977135

95% confidence interval for beta 7 : 0.0038514131536005747 , 0.22711232960438194

95% confidence interval for beta 8 : 0.3926120684225939 , 0.48800680442156474

95% confidence interval for sigma square : 0.36497906775453 ,
0.44325333078843604

The Sum of Squared Errors for the model (SSE) : 327.0238269666074

The Mean Squared error for the model (MSE) : 0.3973558043336663

The Sum of Squares Regression for the model (SSR) : 512.5675209674787

The Total error for the model (SST) : 839.5913479340861

The Model is significant

Multiple Regression Model Summary:

R-squared : 0.6105

Adjusted R-squared : 0.6067

Number of significant terms : 7

Number of non-significant terms: 2


Model significance : significant

SSE on test set: 70.89839076776099

MSE on test set: 7.877598974195665

Chi-square statistic: 2.5648251949055894

P-value: 0.9999970429649789



5 Polynomial regression (Naive)

```
[14]: def multicollinearity_remover(degree,max_corr,X):
    k=np.shape(X)[1] - 1
    c=[]
    for i in range(k+1):
        c.append(i)
    terms=[]
    for x in itertools.combinations_with_replacement(c, degree):
        terms.append(x)
    new_X=np.zeros((np.shape(X)[0],1))
    for j in range(len(terms)):
        e=np.ones((np.shape(X)[0],1))
        for i in terms[j]:
            t=t1=np.reshape(X[:,i],np.shape(e))
            e=np.multiply(e,t)
        new_X=np.concatenate((new_X,e),axis=1)
    new_X=new_X[:,1:]
    t=pd.DataFrame(new_X[:,1:])
    corr_matrix=np.array(t.corr())
    indices_list = []
    adjunct_terms=[]

    # Iterate through the correlation matrix and store indices with value
    ↪ greater than max_corr
    for i in range(corr_matrix.shape[0]):
        for j in range(i + 1, corr_matrix.shape[1]): # Only iterate over upper
    ↪ triangle
            if abs(corr_matrix[i, j]) > max_corr:
                indices_list.append((i, j))

    for i,j in indices_list:
        adjunct_terms.append(terms[j+1])

    utilizable_terms = []
    for element in terms:
        if element not in adjunct_terms:
            utilizable_terms.append(element)
    new_X=np.zeros((np.shape(X)[0],1))
    for j in range(len(utilizable_terms)):
        e=np.ones((np.shape(X)[0],1))
        for i in utilizable_terms[j]:
            t=t1=np.reshape(X[:,i],np.shape(e))
            e=np.multiply(e,t)
        new_X=np.concatenate((new_X,e),axis=1)
    new_X=new_X[:,1:]
```

```
return (new_X,utilizable_terms)
```

```
[15]: def fit_upto_degree(X,Y,degree, test_split = 0.2):  
    X = np.array(X)  
    Y = np.array(Y)  
    n = np.size(Y)  
  
    X_modified,usable_terms = multicollinearity_removal(degree,0.98,X)  
    k = np.shape(X_modified)[1] - 1  
  
    X_train, X_test, Y_train, Y_test = train_test_split(X_modified, Y,  
↳test_size=0.2, random_state=42)  
    beta = fit(X_train,Y_train)  
  
    R2, R2_adj,significance = analyse_model(X_train,Y_train,beta, X_test,Y_test)  
    high_deg_sig=0  
    for i in range(k+1):  
        if ((0 not in usable_terms[i]) and (significance[i]==1)):  
            high_deg_sig=1  
            break  
    plot_residual(X_modified, Y, beta, X_modified.shape[1]-1)  
    return R2, R2_adj,high_deg_sig
```

```
[16]: max_degree = 4  
  
R2 = []  
R2_adj = []  
stop_flag=0  
  
for d in range(1, max_degree + 1):  
    print(f"Degree {d} polynomial regression started")  
    x, y,high_deg_sig = fit_upto_degree(X, Y, d)  
    if(stop_flag==0):  
        if(high_deg_sig==0):  
            print("stop at degree",d-1)  
            stop_flag=1  
        else:  
            print("continue forward selection")  
    if stop_flag:  
        break  
    R2.append(x)  
    R2_adj.append(y)  
    print("\n")  
  
import matplotlib.pyplot as plt  
plt.plot([d for d in range(1, len(R2)+1)], R2, label='R-squared', color='blue')
```

```
plt.plot([d for d in range(1, len(R2)+1)], R2_adj, label='Adjusted R-squared',
        color='red')
plt.xlabel('Polynomial Degree')
plt.ylabel('R-squared')
plt.title('R-squared vs. Polynomial Degree')
plt.legend()
plt.show()
```

Degree 1 polynomial regression started

The Sum of Squared Errors for the model (SSE) : 327.0238269666074

The Mean Squared error for the model (MSE) : 0.3973558043336663

The Sum of Squares Regression for the model (SSR) : 512.5675209674787

The Total error for the model (SST) : 839.5913479340861


The Model is significant

Multiple Regression Model Summary:

```
-----
R-squared           : 0.6105
Adjusted R-squared  : 0.6067
Number of significant terms : 7
Number of non-significant terms: 2
Model significance  : significant
SSE on test set: 70.89839076776099
MSE on test set: 7.877598974195665
-----
```

Chi-square statistic: 2.6207568842880433

P-value: 0.9999964600144139



continue forward selection


Degree 2 polynomial regression started

The Sum of Squared Errors for the model (SSE) : 156.9317787003946

The Mean Squared error for the model (MSE) : 0.19068259866390597
The Sum of Squares Regression for the model (SSR) : 682.6595692336915
The Total error for the model (SST) : 839.5913479340861
The Model is significant
Multiple Regression Model Summary:

R-squared : 0.8131
Adjusted R-squared : 0.8025
Number of significant terms : 29
Number of non-significant terms: 16
Model significance : significant
SSE on test set: 41.07728215659739
MSE on test set: 0.9128284923688309

Chi-square statistic: 3.1198816790193296
P-value: 0.9999851583932261




continue forward selection

Degree 3 polynomial regression started
The Sum of Squared Errors for the model (SSE) : 78.17014791859738
The Mean Squared error for the model (MSE) : 0.09498195372855088
The Sum of Squares Regression for the model (SSR) : 761.4212000154887
The Total error for the model (SST) : 839.5913479340861
The Model is significant
Multiple Regression Model Summary:

R-squared : 0.9069
Adjusted R-squared : 0.8843
Number of significant terms : 32
Number of non-significant terms: 130

Model significance : significant
SSE on test set: 41.83715282127673
MSE on test set: 0.25825402976096745

Chi-square statistic: 4.914237738640554
P-value: 0.9994985823460552




continue forward selection


Degree 4 polynomial regression started
The Sum of Squared Errors for the model (SSE) : 29.958622085983695
The Mean Squared error for the model (MSE) : 0.036401727929506314
The Sum of Squares Regression for the model (SSR) : 809.6327258481024
The Total error for the model (SST) : 839.5913479340861
The Model is significant
Multiple Regression Model Summary:

R-squared : 0.9643
Adjusted R-squared : 0.9173
Number of significant terms : 26
Number of non-significant terms: 443
Model significance : significant
SSE on test set: 1426.7819397618653
MSE on test set: 3.0421789760380924

Chi-square statistic: 0.19699256247957228
P-value: 0.9999999999999998



continue forward selection



6 Principal Component Regression

```
[17]: from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np

def gram_schmidt(vectors):
    num_vectors = vectors.shape[1]
    basis = np.zeros_like(vectors)
    basis[:,0] = vectors[:,0]
    for i in range(num_vectors):
        basis[:, i] = vectors[:, i]
        for j in range(i):
            basis[:, i] -= np.dot(basis[:, j], vectors[:, i]) / np.dot(basis[:, j],
↪basis[:, j]) * basis[:, j]
        basis[:, i] /= np.linalg.norm(basis[:, i])
    return basis

def pcr(X, Y):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
↪random_state=42)
    cov_mat = X_train.T @ X_train
    eig_vals, eig_vecs = np.linalg.eig(cov_mat)
    eig_pairs = [(eig_vals[i], eig_vecs[:,i]) for i in range(len(eig_vals))]
    eig_pairs.sort(key=lambda x: x[0], reverse=True)
    eig_vals_sorted = np.array([X_train[0] for X_train in eig_pairs])
    eig_vecs_sorted = np.array([X_train[1] for X_train in eig_pairs]).T
    p = gram_schmidt(eig_vecs_sorted)
    error = np.linalg.norm(cov_mat - p @ np.diag(eig_vals_sorted) @ p.T)
    if error > 1e-6:
        print("PCR failed: Reconstruction error too large")
        return -1

    total_variance = sum(eig_vals)
    t = 0
    cumulative_sum = 0.0
    for eigenvalue, _ in eig_pairs:
        cumulative_sum += eigenvalue
        t += 1
        if t >= 200:
            break

    selected_eig_pairs = eig_pairs[:t]
    selected_eigenvalues = [pair[0] for pair in selected_eig_pairs]
    selected_eigenvectors = p[:, :t]
    z = X_train @ selected_eigenvectors
    alpha = np.linalg.inv(z.T @ z) @ z.T @ Y_train
    num_features = X_train.shape[1]
```

```

alpha_padded = np.pad(alpha, (0, num_features - len(alpha)),
mode='constant')
coefficients_beta = p @ alpha_padded
Y_train_pred = X_train @ coefficients_beta
SST = np.sum((Y_train - np.mean(Y_train))**2)
SSE = np.sum((Y_train - Y_train_pred)**2)
R_squared = 1 - (SSE / SST)
n = X_train.shape[0]
p = X_train.shape[1] - 1
R_squared_adj = 1 - (1 - R_squared) * ((n - 1) / (n - p - 1))

Y_test_pred = X_test @ coefficients_beta
SSE_test = np.sum((Y_test_pred - Y_test)**2)

print("Mean Squared Error:", SSE/X_train.shape[0])
print("Sum of Squared Errors (SSE):", SSE)
print("Total Sum of Squares (SST):", SST)
print("R-squared (R2):", R_squared)
print("Adjusted R-squared (R2_adjusted):", R_squared_adj)
print("SSE on test set:", SSE_test)
print("MSE on test set:", SSE_test/X_test.shape[0])

return R_squared, R_squared_adj

```

```

[18]: def fit_upto_degree_pcr(X, Y, degree):
    X = np.array(X)
    Y = np.array(Y)
    poly = PolynomialFeatures(degree)
    new_X = poly.fit_transform(X)
    result = pcr(new_X, Y)
    return result

```

```

[19]: max_degree = 4
R2 = []
R2_adj = []
ok = True
X = X_scaled.copy()
for d in range(1, max_degree + 1):
    print(f"Degree {d} PCR started")
    result = fit_upto_degree_pcr(X, Y, d) # Assuming you have a function
    ↪fit_upto_degree_pcr to fit the PCR model
    if result == -1:
        ok = False
        break
    R2.append(result[0])
    R2_adj.append(result[1])
    print("\n")

```

```

if ok:
    plt.plot(np.array([d for d in range(1, max_degree + 1)]), R2,
    ↪label='R-squared', color='blue')
    plt.plot(np.array([d for d in range(1, max_degree + 1)]), R2_adj,
    ↪label='Adjusted R-squared', color='red')
    plt.xlabel('Polynomial Degree')
    plt.ylabel('R-squared')
    plt.title('R-squared vs. Polynomial Degree (PCR)')
    plt.legend()
    plt.show()

```

Degree 1 PCR started

Mean Squared Error: 0.3968735764157856

Sum of Squared Errors (SSE): 327.0238269666073

Total Sum of Squares (SST): 839.5913479340861

R-squared (R2): 0.610496430470266

Adjusted R-squared (R2_adjusted): 0.6066730825485018

SSE on test set: 70.89839076776096

MSE on test set: 0.3441669454745678

Degree 2 PCR started

Mean Squared Error: 0.19045118774319733

Sum of Squared Errors (SSE): 156.9317787003946

Total Sum of Squares (SST): 839.5913479340861

R-squared (R2): 0.8130855217999282

Adjusted R-squared (R2_adjusted): 0.8025280929927354

SSE on test set: 41.07728215659646

MSE on test set: 0.19940428231357504

Degree 3 PCR started

Mean Squared Error: 0.07024401379574416

Sum of Squared Errors (SSE): 57.88106736769319

Total Sum of Squares (SST): 839.5913479340861

R-squared (R2): 0.9310604289692641

Adjusted R-squared (R2_adjusted): 0.9139039955109323

SSE on test set: 29.7723006430607

MSE on test set: 0.1445257312769937

Degree 4 PCR started

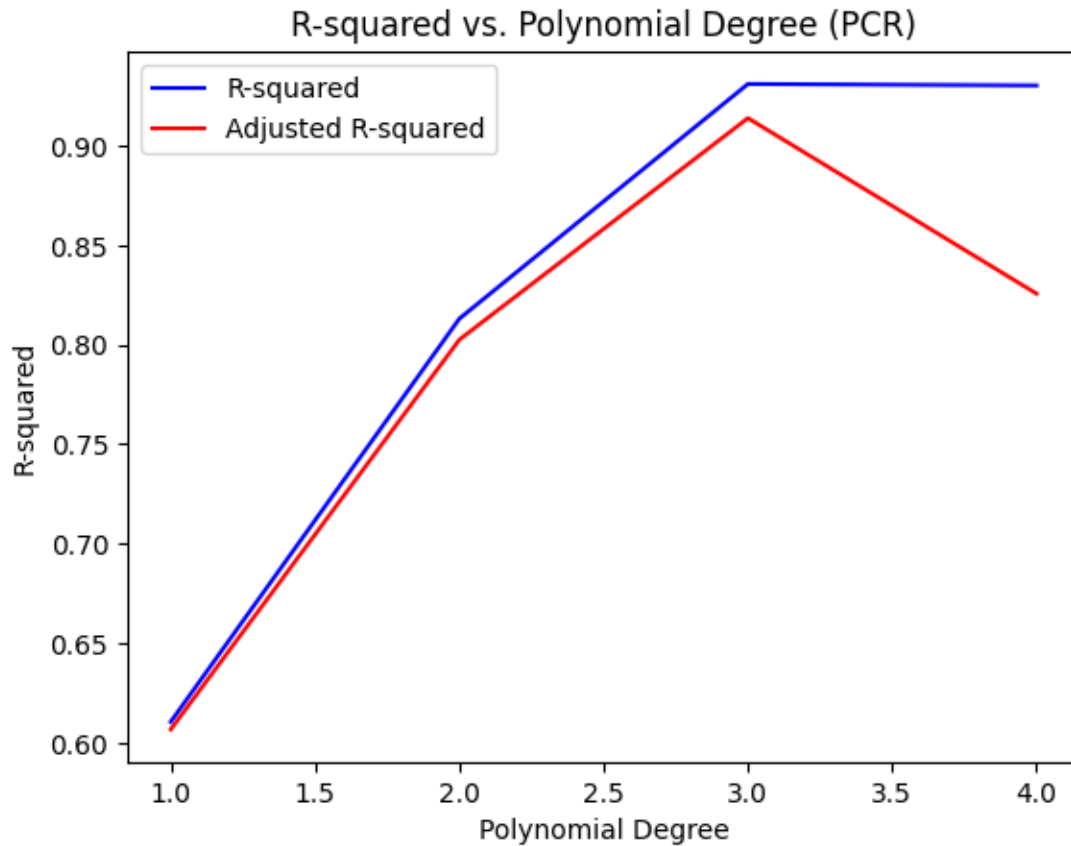
Mean Squared Error: 0.07101037811142759

Sum of Squared Errors (SSE): 58.51255156381634

Total Sum of Squares (SST): 839.5913479340861

R-squared (R2): 0.9303082961636118

Adjusted R-squared (R2_adjusted): 0.825664825965509
SSE on test set: 37.23183790656965
MSE on test set: 0.18073707721635754



7 Orthogonal Polynomial regression

```
[20]: def orthogonal_polynomials(X, degree):  
    q = [np.ones(len(X))]  
    for i in range(1, degree + 1):  
        for j in range(X.shape[1]):  
            q_i = X[:, j]**i - np.sum(q_k * np.dot(X[:, j]**i, q_k) / np.  
dot(q_k, q_k) for q_k in q)  
            q.append(q_i)  
    return np.array(q).T  
  
def orthogonal_poly_regression(X, y, degree):  
    q = orthogonal_polynomials(X, degree)
```

```

        beta = np.linalg.lstsq(q, y, rcond=None)[0]
        return q,beta

max_degree = 5
R2 = []
R2_adj = []

for d in range(1,max_degree+1):
    print(f"Degree {d} orthogonal polynomial regression started")
    X_modified,beta = orthogonal_poly_regression(X, Y, d)
    x,y,_ = analyse_model(X_modified,Y,beta,X_test = None, Y_test = None,
    test_check = False)
    plot_residual(X_modified, Y, beta, X.shape[1]-1)
    R2.append(x)
    R2_adj.append(y)
    print("\n")

import matplotlib.pyplot as plt
plt.plot([d for d in range(1, max_degree + 1)], R2, label='R-squared',
color='blue')
plt.plot([d for d in range(1, max_degree + 1)], R2_adj, label='Adjusted
R-squared', color='red')
plt.xlabel('Polynomial Degree')
plt.ylabel('R-squared')
plt.title('R-squared vs. Polynomial Degree ')
plt.legend()
plt.show()

```

```

Degree 1 orthogonal polynomial regression started
The Sum of Squared Errors for the model ( SSE ) : 396.0277981481325
The Mean Squared error for the model ( MSE ) : 0.3848666648669898
The Sum of Squares Regression for the model ( SSR ) : 633.9722018518676
The Total error for the model ( SST ) : 1030.0
The Model is significant
Multiple Regression Model Summary:

```

```

-----
R-squared           : 0.6155
Adjusted R-squared  : 0.6125
Number of significant terms : 6
Number of non-significant terms: 3
Model significance   : significant
-----

```

```

Chi-square statistic: 2.5584454170794797
P-value: 0.9999971038565955


```

```

<ipython-input-20-9a82c8432e5d>:5: DeprecationWarning: Calling np.sum(generator)
is deprecated, and in the future will give a different result. Use
np.sum(np.fromiter(generator)) or the python sum builtin instead.

```

```
q_i = X[:, j]**i - np.sum(q_k * np.dot(X[:, j]**i, q_k) / np.dot(q_k, q_k) for
q_k in q)
```



Degree 2 orthogonal polynomial regression started

The Sum of Squared Errors for the model (SSE) : 232.5395203884522

The Mean Squared error for the model (MSE) : 0.22598592846302448

The Sum of Squares Regression for the model (SSR) : 797.4604796115478

The Total error for the model (SST) : 1030.0


The Model is significant

Multiple Regression Model Summary:

```
-----
R-squared           : 0.7742
Adjusted R-squared  : 0.7707
Number of significant terms : 13
Number of non-significant terms: 4
Model significance   : significant
-----
```

Chi-square statistic: 3.2729750608367136

P-value: 0.9999781480839348



Degree 3 orthogonal polynomial regression started

The Sum of Squared Errors for the model (SSE) : 164.73741352455346

The Mean Squared error for the model (MSE) : 0.1600946681482541

The Sum of Squares Regression for the model (SSR) : 865.2625864754466

The Total error for the model (SST) : 1030.0


The Model is significant

Multiple Regression Model Summary:

```
-----
R-squared           : 0.8401
Adjusted R-squared  : 0.8362
Number of significant terms : 18
Number of non-significant terms: 7
Model significance   : significant
-----
```

Chi-square statistic: 3.3308788564490452


P-value: 0.9999748440700084



Degree 4 orthogonal polynomial regression started
The Sum of Squared Errors for the model (SSE) : 132.12062533564142
The Mean Squared error for the model (MSE) : 0.12839710916971955
The Sum of Squares Regression for the model (SSR) : 897.8793746643586
The Total error for the model (SST) : 1030.0
The Model is significant
Multiple Regression Model Summary:

R-squared : 0.8717
Adjusted R-squared : 0.8676
Number of significant terms : 24
Number of non-significant terms: 9
Model significance : significant

Chi-square statistic: 3.5736295162419576
P-value: 0.9999559501194862





Degree 5 orthogonal polynomial regression started
The Sum of Squared Errors for the model (SSE) : 125.946129569449
The Mean Squared error for the model (MSE) : 0.12239662737555783
The Sum of Squares Regression for the model (SSR) : 904.053870430551
The Total error for the model (SST) : 1030.0
The Model is significant
Multiple Regression Model Summary:

R-squared : 0.8777

Adjusted R-squared : 0.8728
Number of significant terms : 28
Number of non-significant terms: 13
Model significance : significant

Chi-square statistic: 3.396723647707346
P-value: 0.999970578231546





8 LASSO

```
[21]: # from sklearn.linear_model import Lasso

# def fit_upto_degree_lasso(X, Y, degree, alpha=10):
#     poly = PolynomialFeatures(degree)
#     new_X = poly.fit_transform(X)
#     # print(new_X, Y)
#     lasso_model = Lasso(alpha=alpha, fit_intercept = False)
#     lasso_model.fit(new_X, Y)
#     beta = lasso_model.coef_
#     R2, R2_adj, temp = analyse_model(new_X, Y, beta)
#     return R2, R2_adj
```

```
[22]: # max_degree = 10
# R2 = []
# R2_adj = []
# for d in range(1, max_degree + 1):
#     r2, r2_adj = fit_upto_degree_lasso(X, Y, d)
```

```

#     R2.append(r2)
#     R2_adj.append(r2_adj)

# import matplotlib.pyplot as plt
# plt.plot([d for d in range(1, max_degree + 1)], R2, label='R-squared',
# ↪color='blue')
# plt.plot([d for d in range(1, max_degree + 1)], R2_adj, label='Adjusted
# ↪R-squared', color='red')
# plt.xlabel('Polynomial Degree')
# plt.ylabel('R-squared')
# plt.title('R-squared vs. Polynomial Degree')
# plt.legend()
# plt.show()

```

9 RIDGE

```

[23]: from sklearn.linear_model import Ridge

def fit_upto_degree_ridge(X, Y, degree, alpha=100):
    poly = PolynomialFeatures(degree)
    new_X = poly.fit_transform(X)

    X_train, X_test, Y_train, Y_test = train_test_split(new_X, Y, test_size=0.
↪2, random_state=42)
    ridge_model = Ridge(alpha=alpha, fit_intercept=False)
    ridge_model.fit(X_train, Y_train)

    beta = ridge_model.coef_
    R2, R2_adj, temp = analyse_model(X_train, Y_train, beta, X_test, Y_test)

    return R2, R2_adj

```

```

[24]: max_degree = 4
R2 = []
R2_adj = []

for d in range(1, max_degree + 1):
    print(f"Degree {d} Ridge regression started")
    r2, r2_adj = fit_upto_degree_ridge(X, Y, d)
    R2.append(r2)
    R2_adj.append(r2_adj)
    print("\n")

import matplotlib.pyplot as plt

```

```
plt.plot([d for d in range(1, max_degree + 1)], R2, label='R-squared',
        color='blue')
plt.plot([d for d in range(1, max_degree + 1)], R2_adj, label='Adjusted
        R-squared', color='red')
plt.xlabel('Polynomial Degree')
plt.ylabel('R-squared')
plt.title('R-squared vs. Polynomial Degree')
plt.legend()
plt.show()
```

Degree 1 Ridge regression started

The Sum of Squared Errors for the model (SSE) : 344.79013502529597

The Mean Squared error for the model (MSE) : 0.4189430559238104

The Sum of Squares Regression for the model (SSR) : 494.80121290879015

The Total error for the model (SST) : 839.5913479340861

The Model is significant

Multiple Regression Model Summary:

```
-----
R-squared           : 0.5893
Adjusted R-squared  : 0.5853
Number of significant terms : 5
Number of non-significant terms: 4
Model significance   : significant
SSE on test set: 75.1199902081647
MSE on test set: 8.346665578684966
-----
```

Degree 2 Ridge regression started

The Sum of Squared Errors for the model (SSE) : 215.67440559227094

The Mean Squared error for the model (MSE) : 0.2620588160294908

The Sum of Squares Regression for the model (SSR) : 623.9169423418152

The Total error for the model (SST) : 839.5913479340861

The Model is significant

Multiple Regression Model Summary:

```
-----
R-squared           : 0.7431
Adjusted R-squared  : 0.7286
Number of significant terms : 8
Number of non-significant terms: 37
Model significance   : significant
SSE on test set: 51.81956868639005
MSE on test set: 1.1515459708086677
-----
```


Degree 3 Ridge regression started

The Sum of Squared Errors for the model (SSE) : 132.01464410647222
The Mean Squared error for the model (MSE) : 0.16040661495318617
The Sum of Squares Regression for the model (SSR) : 707.5767038276139
The Total error for the model (SST) : 839.5913479340861
The Model is significant
Multiple Regression Model Summary:

R-squared : 0.8428
Adjusted R-squared : 0.8036
Number of significant terms : 2
Number of non-significant terms: 163
Model significance : significant
SSE on test set: 44.40149619081102
MSE on test set: 0.26909997691400617

Degree 4 Ridge regression started
The Sum of Squared Errors for the model (SSE) : 80.85462098685512
The Mean Squared error for the model (MSE) : 0.09824376790626382
The Sum of Squares Regression for the model (SSR) : 758.736726947231
The Total error for the model (SST) : 839.5913479340861
The Model is significant
Multiple Regression Model Summary:

R-squared : 0.9037
Adjusted R-squared : 0.7591
Number of significant terms : 0
Number of non-significant terms: 495
Model significance : significant
SSE on test set: 34.839366553925906
MSE on test set: 0.07038255869479981



[24] :